

# Implementing variant visitation using lambdas

**Bloomberg**

**Vittorio Romeo**

<https://vittorioromeo.info>  
vittorio.romeo@outlook.com

**C++::London**

11/05/2017

# About me

- Developer at **Bloomberg L.P.**
- Modern C++ enthusiast
  - Conference talks
  - Video tutorials & articles
  - Open-source projects

# Overview

1. What is a *variant* and why is it useful?
2. Variant visitation

part 1

## What is a *variant* and why is it useful?

- Variant types and use cases
- Variant implementations
- `std::variant`

To understand variants, let's begin by understanding `struct` and `enum class`.

What is a **struct** ?

A `struct` models aggregation of types.

```
struct point
{
    int _x;
    int _y;
};
```

A `point` is an `int` **AND** an `int` .

What is an **enum class** ?



An `enum class` models a choice between values.

```
enum class traffic_light  
{  
    red,  
    yellow,  
    green  
};
```

A `traffic_light` is **EITHER** `red` **OR** `yellow` **OR** `green` .

What is a **variant** ?

A `variant` models a choice between types.

```
struct on { int _temperature; };  
struct off { };  
  
using oven_state = std::variant<on, off>;
```

- The oven is `off` .

*...or...*

- The oven is `on` , with `_temperature` = 200°C.

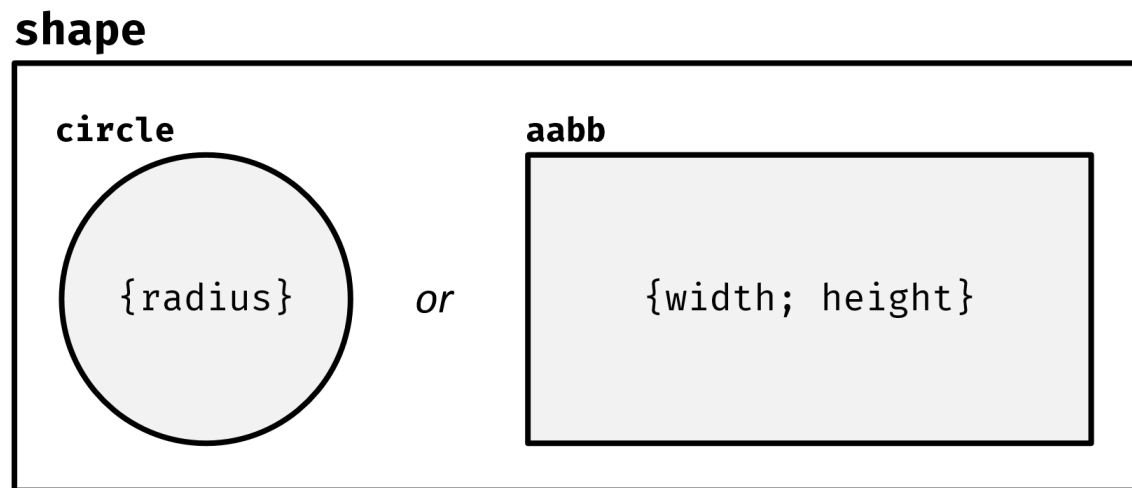
	<b>struct</b>	<b>enum class</b>	<b>variant</b>
<b>model</b>	<i>aggregation: types</i>	<i>choice: values</i>	<i>choice: types</i>
<b>class</b>	product type	sum type	sum type

Variant types can be thought of as **type-safe tagged unions** that:

- Require significantly less boilerplate.
- Automatically deal with constructors/destructors and assignment.
- Immensely increase **safety**.

Those claims are easy to verify through an *example*. Let's...

- Define a `shape` **sum type** which can be a **circle** or an **axis-aligned bounding box**.
- Define an `area` function that, given a `shape` instance, calculates and returns its **surface area**.



We will start with a "*traditional*" *tagged union* first.

```
enum class shape_type{circle, aabb};

struct shape
{
    shape_type _type;

    union
    {
        struct { float _radius;          } circle_data;
        struct { float _width, _height; } aabb_data;
    } _u;
};
```



When defining a "*traditional*" *tagged union* type, we need:

- An `enum class` with a value for every *alternative*.
- A `struct` for every alternative that defines its *state*.
- An `union` that wraps all the state `struct` types.
- An interface `struct` that wraps the `union` and `enum class`.

```
auto area(const shape& s)
{
    switch(s._type)
    {
        case shape_type::circle:
        {
            const auto& r = s._u.circle_data._radius;
            return pi * r * r;
        }
        case shape_type::aabb:
        {
            const auto& x = s._u.aabb_data;
            return x._width * x._height;
        }
    };

    assert(false);
    __builtin_unreachable();
}
```

[on [gcc.godbolt.org](https://gcc.gnu.org/godbolt.org)]

When visiting it, we need:

- A `switch` with a `case` for every alternative.
- Every `case` needs to access the relevant state by going through the `union`.
- An *unreachable* guard at the end of the visitation.

Let's now see how `std::variant` compares.

```
struct circle { float _radius; };  
struct aabb { float _width, _height; };  
using shape = std::variant<circle, aabb>;
```

```
auto area(const shape& s)  
{  
    struct {  
        auto operator()(const circle& x) const  
        {  
            return pi * x._radius * x._radius;  
        }  
        auto operator()(const aabb& x) const  
        {  
            return x._width * x._height;  
        }  
    } visitor;  
  
    return std::visit(visitor, s);  
}
```

[on [gcc.godbolt.org](https://gcc.godbolt.org)]

- Defining a variant type is trivial: all that's required is listing the alternative types.
- Visitation allows easy access to the members of currently active variant alternative.
- `visitor` is checked to be *exhaustive* at compile-time.

Examples of structures elegantly modeled by variants are:

- JSON
- ASTs (*abstract syntax trees*)
- State machines
- Error handling

```
using json_value = std::variant<
    json_object,
    json_array,
    json_string,
    json_number,
    json_boolean,
    json_null
>;
```

```
template <typename T>
json_value serialize(const T&);

template <typename T>
T deserialize(const json_value&);
```



```
struct success          { data _contents; }  
struct file_not_found   { };  
struct invalid_permissions { };  
  
using file_open_result = std::variant  
<  
    success,  
    file_not_found,  
    invalid_permissions  
>;
```

```
file_open_result open_file(const std::string& path);
```

The examples used `std::variant`, standardized in C++17.

Several other implementations are available today:

- `boost::variant`
- `eggs::variant`
- `type_safe::variant`
- `bdlb::Variant`

They differ in their:

- Interface
- Default initialization semantics
- Existence of an "empty" state
- Strategy of dealing with exceptions
- Rules for duplicate types

The visitation techniques that will be covered in this talk will be applicable to **any variant implementation**.

For simplicity, we're going to keep using `std::variant` for the rest of the talk.

`std::variant<Ts ... >` can be constructed/assigned with any instance of `Ts ...` or with any other `variant` instance.

```
std::variant<int, std::string> v0{1};  
v0 = "hello!"s;  
  
std::variant<int, std::string> v1{"bye!"s};  
v0 = std::move(v1);
```

The active alternative in an `std::variant` instance can be accessed with any of the following:

- `std::variant::get`
- `std::variant::get_if`

```
std::variant<int, std::string> v0{1};

assert(v0.holds_alternative<int>());
assert(v0.get<int>() == 1);

auto v0_str = v0.get_if<std::string>();
if(v0_str != nullptr)
{
    // ...
}
```

- `get<T>` requires the user to be aware of the currently active alternative of the variant. In case of error, an *exception* will be thrown.
- `get_if<T>` requires syntactical overhead that does not scale well with a large number of types.

That's why the standard library provides `std::visit`, which allows **variant visitation**. We'll take a look at it in **part 2**.

## part 1 - recap

- Variants are **sum types** that model a **choice between types**.
- Variants are **safer** and **more expressive** than "*traditional tagged unions*".
- *Recursive node-based structures, error handling, state machines, etc...* are examples of concepts that can be **elegantly** modeled with variants.
- `std::variant` is a **vocabulary type** introduced in C++17.



part 2

## Variant visitation

- What is "*visitation*"?
- "Traditional" visitation
- "Lambda-based" visitation
  - Arbitrary *function object* overloading

## What is "*visitation*"?

Visitation can be defined as an **abstraction** over accessing the currently active variant *alternative* in an **exhaustive** and **expressive** manner.

# "Traditional" visitation

## overview

- Visitation requires a `Callable` object which can be invoked with every possible variant alternative.
- The "traditional" way of creating such as object is defining a `struct`.

## "Traditional" visitation

example - one variant

```
struct printer
{
    void operator()(int x)      { cout << x << "i\n"; }
    void operator()(float x)   { cout << x << "f\n"; }
    void operator()(double x) { cout << x << "d\n"; }
};
```

```
using my_variant = std::variant<int, float, double>;
my_variant v0{20.f};

// Prints "20f".
std::visit(printer{}, v0);
```

[on [gcc.godbolt.org](https://gcc.godbolt.org)]

## "Traditional" visitation

example - two variants

```
struct collision_resolver
{
    void operator()(circle, circle) { /* ... */ }
    void operator()(circle, aabb)   { /* ... */ }
    void operator()(aabb, circle)  { /* ... */ }
    void operator()(aabb, aabb)    { /* ... */ }
};
```

```
using my_variant = std::variant<circle, aabb>;
my_variant v0{circle{}};
my_variant v1{aabb{}};

std::visit(collision_resolver{}, v0, v1);
```

[on [gcc.godbolt.org](https://gcc.gnu.org/godbolt.org)]

## "Traditional" visitation

### shortcomings

- **Syntactical overhead:** a `struct` with multiple `operator( )` overloads must be defined.
- **Lack of locality:** sometimes the `struct` cannot be defined locally (*e.g. contains template methods*).
- **Readability impact:** the visitation logic is defined far away from the visitation site.

| Can we do better?

**Let's take some inspiration from *Rust*.**

## Rust *variants*: **enum**

example

```
enum MyVariant {  
    IntTag(i32),  
    FloatTag(f32),  
    DoubleTag(f64)  
}
```

```
let v0 = FloatTag(2.0);  
match v0 {  
    IntTag(x)      ⇒ println!("{}", x),  
    FloatTag(x)    ⇒ println!("{}", x),  
    DoubleTag(x)   ⇒ println!("{}", x)  
}
```

[on [rust.godbolt.org](https://rust.godbolt.org)]



Rust has *language-level* **variants** and **pattern-matching**. This allows expressive and safe local visitation of variants.

```
match v0 {  
    IntTag(x)    ⇒ println!("{}", x),  
    // ^^^^^^^^^  
    // pattern  
  
    FloatTag(x) ⇒ println!("{}", x),  
    // ^^^  
    // decomposition  
  
    DoubleTag(x) ⇒ println!("{}", x)  
    // ^^^^^^^^^^^^^^^^^^^^^^^^^  
    // expression  
}
```

C++

```
struct printer
{
    void operator()(int x)      { cout << x << "i\n"; }
    void operator()(float x)   { cout << x << "f\n"; }
    void operator()(double x) { cout << x << "d\n"; }
};

my_variant v0{20.f};
std::visit(printer{}, v0);
```

Rust

```
let v0 = FloatTag(2.0);
match v0 {
    IntTag(x)      ⇒ println!("{}i", x),
    FloatTag(x)    ⇒ println!("{}f", x),
    DoubleTag(x)   ⇒ println!("{}d", x)
}
```

# "Lambda-based" visitation

## overview

- Visitation is done **locally** by using a set of exhaustive *lambda expressions*.
- *"Pattern matching"-like* syntax.
- Minimal syntactical boilerplate.
- No additional run-time overhead.

## "Lambda-based" visitation

example

```
using my_variant = std::variant<int, float, double>;  
my_variant v0{20.f};  
  
// Prints "20f".  
match(v0)([](int x) { cout << x << "i\n"; },  
          [](float x) { cout << x << "f\n"; },  
          [](double x) { cout << x << "d\n"; });
```

Almost no syntactical overhead!

Let's compare again to Rust's `match` ...

C++

```
my_variant v0{20.f};

// Prints "20f".
match(v0)([](int x)    { cout << x << "i\n"; },
          [](float x)  { cout << x << "f\n"; },
          [](double x){ cout << x << "d\n"; });
```

Rust

```
let v0 = FloatTag(2.0);
match v0 {
    IntTag(x)    => println!("{}", x),
    FloatTag(x)  => println!("{}", x),
    DoubleTag(x) => println!("{}", x)
}
```

(Note: Rust's `match` can be *way* more powerful than shown here.)

Before diving into `match`'s implementation, let's look at an additional example.

## "Lambda-based" visitation

example

```
struct payload { int _data; };  
struct error    { std::string _what; };  
  
using request = std::variant<payload, error>;
```

```
match(make_request("legitwebsite.com"))(  
    [](const payload& p)  
    {  
        match(make_request("abc.com", p._data))(  
            [](const payload& p){ /* ... */ },  
            [](const error& e)  { /* ... */ });  
    },  
    [](const error& e){ /* ... */ });
```

# "Lambda-based" visitation

## implementation - overview

- `match` will be a function that takes  $N_v$  *variants* and returns a function that takes  $N_f$  *function objects*.

```
match(v0, v1, ... , vN_v)(f0, f1, ... , fN_f);
```

- In order to create a *visitor* from the passed function objects, an *overload set* must be built out of them.
- Internally, `std::visit` will be called with the *variants* and the newly-built *overload set*.



Let's begin with the creation of an **overload set**.

Let's start by looking at *non-member functions*.

```
int foo(float) { return 0; }  
int foo(char)  { return 1; }
```

`foo` is an *overloaded function*.

```
auto x0 = foo(0.f); // `x0` is `0`.  
auto x1 = foo('a'); // `x1` is `1`.
```

| Can we "*generate*" an overload set out of them?

There is no way of defining *overloaded functions* locally:

```
void f()  
{  
    int foo(float) { return 0; }  
    int foo(char)  { return 1; }  
    // Nope. COMPILER ERROR!  
}
```

...this doesn't help.

| What about *function objects*?

```
struct foo
{
    int operator()(float) { return 0; }
    int operator()(char)  { return 1; }
};
```

`foo` is a *function object*.

```
auto x0 = foo{}(0.f); // `x0` is `0`.
auto x1 = foo{}('a'); // `x1` is `1`.
```

| Can we "*generate*" an overload set out of them?

We can define them locally:

```
void f()  
{  
    struct foo  
    {  
        int operator()(float) { return 0; }  
        int operator()(char)  { return 1; }  
    };  
}
```



We can also compose them via inheritance!

```
struct foo_float { int operator()(float) { return 0; } };  
struct foo_char { int operator()(char) { return 1; } };
```

```
struct foo : foo_float, foo_char  
{  
    using foo_float::operator();  
    using foo_char::operator();  
};
```

```
auto x0 = foo{}(0.f); // `x0` is `0`.  
auto x1 = foo{}('a'); // `x1` is `1`.
```

| "Why do we need the `using` declarations?"

```
struct foo : foo_float, foo_char
{
    // using foo_float::operator();
    // using foo_char::operator();
};

foo{}(0.f); // Nope. COMPILER ERROR!
```

[on [gcc.godbolt.org](https://gcc.gnu.org/bugzilla/show_bug.cgi?id=84646)]

Without the `using-declarations` the code would generate a compiler error.

The reason is that the call to `foo::operator()` would be **ambiguous** because *name resolution* is performed before *overload resolution*. See [§13.2 "Member name lookup"](#).

We're on the right track - all that's needed is some *generalization* via `template` metaprogramming.

First attempt:

```
template <typename ... TFs>
struct overload_set : TFs ...
{
    using TFs :: operator() ... ;
};
```

```
using foo = overload_set<foo_float, foo_char>;
auto x0 = foo{}(0.f);
auto x1 = foo{}('a');
```

...will it work?

The following code...

```
template <typename ... TFs>
struct overload_set : TFs ...
{
    using TFs :: operator() ... ;
};
```

...works perfectly in C++17, thanks to [P0195](#) (*"Pack expansions in using-declarations"*).

It unfortunately causes a compiler error in C++14:

```
error: parameter packs not expanded with '...':  
struct overload_set : TFs ...  
{ using TFs::operator() ... ; };  
//          ^^^^
```

This happens because C++14 *using*-declarations don't support multiple comma-separated items. E.g.

```
struct a : b, c  
{  
    using b::foo, c::bar; // Error before C++17.  
};
```

The problems therefore are:

- Only one of the base classes' `operator()` can be introduced in the scope of `overload_set` with the `using` keyword.
- There's no way of generating all the *using-declarations* at once.

**...how can we work around these limitations?**

The first one gives us an hint: let's try to separately match the first base class.



```
template <typename TF, typename ... TFs>
struct overload_set : TF, TFs ...
{
    using TF::operator();
    // ... what to do with `TFs ...`?
};
```

Can you see the pattern? This looks like a job for **recursion**!

Recursive case:

```
template <typename TF, typename ... TFs>
struct overload_set : TF, overload_set<TFs ... >
{
    using TF::operator();
    using overload_set<TFs ... >::operator();
};
```

Base case (*specialization*):

```
template <typename TF>
struct overload_set<TF> : TF
{
    using TF::operator();
};
```

Great! Let's try out our new `overload_set` :

```
struct foo_float
{
    constexpr int operator((float) { return 0; }
};

struct foo_char
{
    constexpr int operator((char) { return 1; }
};
```

```
overload_set<foo_float, foo_char> o;

static_assert(o(0.f) == 0);
static_assert(o('a') == 1);
```

[\[on gcc.godbolt.org\]](https://gcc.godbolt.org)

In order to support *lambda expressions* and automatically deduce the passed *function object* types, we need to make two improvements:

- Have a *perfect-forwarding* constructor inside `overload_set`.
  - Lambdas are **not** `DefaultConstructible` !
- Provide an `overload` interface *variadic function template* that deduces the types of *function objects* and returns an `overload_set`.

The *base case's perfect-forwarding* constructor is trivial to implement:

```
template <typename TF>
struct overload_set<TF> : TF
{
    using TF::operator();

    template <typename TFFwd>
    overload_set(TFFwd&& f)
        : TF{std::forward<TFFwd>(f)}
    {
    }
};
```

Note that we're not using `TF` as the type of the `f` argument because we want a *forwarding reference*, not an *rvalue reference*.

The *recursive case's* constructor is slightly more complicated:

```
template <typename TF, typename ... TFs>
struct overload_set : TF, overload_set<TFs ... >
{
    using TF::operator();
    using overload_set<TFs ... >::operator();

    template <typename TFFwd, typename ... TRest>
    overload_set(TFFwd&& f, TRest&& ... rest) :

        // Construct the current function.
        TF{std::forward<TFFwd>(f)},

        // Construct the base class.
        // Pay attention to the ellipses' positions.
        overload_set<TFs ... >{
            std::forward<TRest>(rest) ... }

    { }
};
```

The last piece of the puzzle is the aforementioned `overload` *variadic function template* interface:

```
template <typename ... TFs>
auto overload(TFs&& ... fs)
{
    return
        overload_set<std::remove_reference_t<TFs> ... >(
            std::forward<TFs>(fs) ... );
}
```

`std::remove_reference_t` is used to convert `T&` to `T` for the special *lvalue forwarding-reference* template deduction rules.

With everything in place, we can finally write the code below:

```
auto o = overload([](float){ return 0; },
                  [] (char) { return 1; });

static_assert(o(0.f) == 0);
static_assert(o('a') == 1);
```

[on [gcc.godbolt.org](https://gcc.godbolt.org)]

Note that lambdas in C++17 are *implicitly* `constexpr` if possible - `static_assert` therefore works with them. (See [N4487](#) and [P0170](#).)



Now that we have a working `overload_set`, let's go back to the original task: making the snippet below compile and work.

```
using my_variant = std::variant<int, float, double>;  
my_variant v0{20.f};  
  
// Prints "20f".  
match(v0)([](int x)    { cout << x << "i\n"; },  
          [](float x) { cout << x << "f\n"; },  
          [](double x){ cout << x << "d\n"; });
```

Let's implement `match` !

Remember:

`match` will be a function that takes  $N_v$  *variants* and returns a function that takes  $N_f$  *function objects*.

```
match(v0, v1, ... , vN_v)(f0, f1, ... , fN_f);
```

Therefore it will roughly look like this:

```
template <typename ... TVariants>
auto match(TVariants&& ... vs)
{
    return [](auto&& ... fs){ /* ... */ };
}
```

Thanks to `overload_set` and `std::visit`, the real implementation is *trivial*.

```
template <typename ... TVariants>
constexpr auto match(TVariants&& ... vs)
{
    return [&vs ... ](auto&& ... fs) → decltype(auto)
    {
        auto visitor =
            overload(std::forward<decltype(fs)>(fs) ... );

        return std::visit(visitor,
            std::forward<TVariants>(vs) ... );
    };
}
```

Finally...

```
using my_variant = std::variant<int, float, double>;  
my_variant v0{20.f};
```

```
// Prints "20f".
```

```
match(v0)([](int x)    { cout << x << "i\n"; },  
          [](float x)  { cout << x << "f\n"; },  
          [](double x){ cout << x << "d\n"; });
```

[on [gcc.godbolt.org](https://gcc.godbolt.org)] | [on [melpon.org/wandbox](https://melpon.org/wandbox)]

...compiles and prints "20f" !

Here's an example with *multiple variants*:

```
struct circle { /* ... */ };
struct aabb    { /* ... */ };

using shape = std::variant<circle, aabb>;
shape s0{circle{}};
shape s1{aabb{}};

match(s0, s1)(
    [](circle, circle){ cout << "circle vs circle\n"; },
    [](circle, aabb)   { cout << "circle vs aabb\n"; },
    [](aabb,   circle){ cout << "aabb vs circle\n"; },
    [](aabb,   aabb)   { cout << "aabb vs aabb\n"; });
```

[on [gcc.godbolt.org](https://gcc.godbolt.org)] | [on [melpon.org/wandbox](https://melpon.org/wandbox)]

One more, with *nested variants*:

```
struct format { /* ... */ };
struct timeout { /* ... */ };
using error = std::variant<format, timeout>;

struct accept { /* ... */ };
struct reject { /* ... */ };
using ok = std::variant<accept, reject>;
```

```
using response = std::variant<error, ok>;
response r{error{timeout{}}};

match(r)(
    [](ok x)    { match(x)([](accept) { /* ... */ },
                          [](reject) { /* ... */ }); },

    [](error x){ match(x)([](format) { /* ... */ },
                          [](timeout){ /* ... */ }); });
```

[on [gcc.godbolt.org](https://gcc.godbolt.org)] | [on [melpon.org/wandbox](https://melpon.org/wandbox)]

## part 2 - recap

- Variant visitation works by invoking the *overload* matching the active *alternative* on an overloaded visitor *function object*.
- `using`-declarations and inheritance can be used to **create overloads** of arbitrary *function objects*.
- Lambdas are just syntactic sugar for *function objects*.
- Overloads of lambdas built *on-the-spot* **reduce boilerplate** and **increase locality** when visiting variants.

## part 2 - recap

- Note that `struct`-based *"traditional"* visitation is **still useful!**
- `struct` visitors can...
  - Contain internal state and a richer interface.
  - Be more easily shared/reused in a large codebase.
  - Provide *customization points* (e.g. via a `template` parameter).
- If your *visitation logic* is simple and local, prefer *"lambda-based"* visitation.



**match** can be implemented in less than  
*40 well-formatted* lines of C++14 code

"...what's the catch?"

The image shows a YouTube video player. The video frame displays a man at a podium with an 'accu 2017' sign. A code snippet is overlaid on the right side of the video frame. Below the video frame, the video title, channel information, and engagement metrics are visible.

accu 2017

Example:

```
expr e{make_unique<r_expr>(1, minus{},
    make_unique<r_expr>(3, plus{}, 7))};

std::cout << match_recursively<int>(e)(
    [](auto, number x) { return x; },
    [](auto recurse, const std::unique_ptr<r_expr>& x)
    {
        const auto& [lhs, op, rhs] = *x;
        return match(op)(
            [&](plus) { return lhs + recurse(rhs); },
            [&](minus){ return lhs - recurse(rhs); });
    });
```

[on gcc.godbolt.org] | [on melpon.org/wandbox]

1:27:24 / 1:30:55 vittorioromeo.info | vittorio.romeo@outlook.com | vromeo5@bloomberg.net | @supahvee1234

Implementing variant Visitation Using Lambdas - Vittorio Romeo [ACCU 2017]

ACCU Conference

Subscribed 1,519

588 views

+ Add to Share ... More

10 0

<https://www.youtube.com/watch?v=mqei4JJRQ7s>

# Thanks!

- [vromeo5@bloomberg.net](mailto:vromeo5@bloomberg.net)
- <https://vittorioromeo.info>
- [vittorio.romeo@outlook.com](mailto:vittorio.romeo@outlook.com)
- [@supahvee1234](#)